

Design/Implementation Notes

1. Error mapping Per Function :

- `_owner_repo()` :
 1. Raises 500 if GITHUB_OWNER or GITHUB_REPO missing.
- `_auth_headers()`
 1. Raises 401 if GITHUB_TOKEN missing.
- `create_issue()`
 1. Delegates rate checks to `map_rate_limit()`.
 2. GitHub 401/403 -> mapped to 401 (detail = r.text).
 3. Any other `r.status_code >= 400` -> mapped to 400 (detail = r.text).
 4. On success -> returns `r.json()`.
- `list_issues()`
 1. Delegates rate checks to `map_rate_limit()`.
 2. If ETag returned and server replies 304 -> raises 304 ("cached").
 3. GitHub 401/403 -> mapped to 401.
 4. GitHub 404 -> mapped to 404 ("not found").
 5. Any other `r.status_code >= 400` -> mapped to 400.
 6. On success -> returns (`r.json()`, headers).
- `get_issue()`
 1. Delegates rate checks to `map_rate_limit()`.
 2. GitHub 404 -> mapped to 404 ("not found").
 3. GitHub 401/403 -> mapped to 401.
 4. Any other `r.status_code >= 400` -> mapped to 400.
 5. On success -> returns `r.json()`.
- `update_issue()`
 1. Same mapping as `get_issue` (404 -> 404, 401/403 -> 401, others -> 400).
- `create_comment()`
 1. Same mapping as `get_issue/update_issue` (404 -> 404, 401/403 -> 401, others -> 400).

2. Pagination strategy :

- `list_issues(state, labels, page, per_page)`
 1. Builds params = {"state": state, "page": page, "per_page": per_page} and adds labels if provided.
 2. Uses headers by calling `_auth_headers()` (includes Authorization and Accept).
 3. If a cached ETag exists for the key, sets If-None-Match header.

Design/Implementation Notes

4. ETag caching :
 - a. In-process dict `_etag_cache`: Dict[Tuple[str, str, int, int], str]
 - b. Cache key = (state, labels or "", page, per_page)
 - c. On response, if `r.headers` contains "etag" it stores `_etag_cache[key] = etag`.
5. Response handling
 - a. Calls `map_rate_limit(r.headers)` and raises its exception if present.
 - b. If `r.status_code == 304` -> raises `HTTPException(status_code=304, detail="cached")`.
 - c. If `r.status_code >= 400` -> handles 401/403 -> 401, 404 -> 404, else -> 400 (via `HTTPException`).
 - d. On success -> returns (`r.json()`, `dict(r.headers)`) — body plus raw GitHub response headers (so callers can read `Link`, `X-RateLimit-*`, etc.).
6. Transport
 - a. Uses `httpx.AsyncClient(timeout=20)` for the request.
7. Implementation facts
 - a. Pagination is page-based and forwarded directly to GitHub.
 - b. ETag caching is per-process, keyed by (state, labels, page, per_page).
 - c. The function returns raw headers alongside the JSON payload.

3. Webhook dedupe :

- Signature verification
 1. Verify `X-Hub-Signature-256` (HMAC-SHA256) against `WEBHOOK_SECRET` before processing.
- Idempotency key
 1. Use `X-GitHub-Delivery` header as the canonical delivery id (fallback: payload hash).
- File-backed index check
 1. Load `events_index.json` (or similar index file) and check if delivery id is present.
 2. If present -> treat as duplicate and return 200 immediately.
- Atomic reservation + append
 1. If not present -> atomically mark delivery id as "processing" in the index (write `tmp` + `os.replace` or use file lock).
 2. Append full event (or minimal metadata) to `events.jsonl` (append-only line-delimited JSON).
 3. Update index entry to "processed" with timestamp/offset and persist atomically.
- Failure semantics
 1. On duplicate: 200 (idempotent).
 2. On transient processing failure: return 5xx so GitHub may retry.

Design/Implementation Notes

4. Security trade-offs :

- File-backed events_index.json with no strong locking/GC
 1. Risk: concurrent writes may cause index corruption or duplicate processing (integrity issue more than confidentiality).
- .env secrets on disk for dev
 1. Risk: accidental commit or backup exposure of GITHUB_TOKEN / WEBHOOK_SECRET.
- Returning upstream r.text in HTTPException details
 1. Risk: may leak GitHub error bodies or internal info to clients.